Protected Legal Thought in the Age of AI

Work Product and the Civil Rights of Legal Preparation

Author Name Redacted for Blind Review

Abstract

Federal courts have begun dividing on whether AI-assisted legal preparation is protected work product. In *United States v. Heppner*, a court held that a defendant's AI-generated legal analysis was discoverable because it was not created under counsel's direction. On the same day, in *Warner v. Gilbarco, Inc.*, a court held that a pro se litigant's AI-assisted materials were protected because they reflected her own litigation strategy. Together, these decisions expose a growing fracture: materially similar legal preparation receives different protection based not on the nature of the work, but on the status of the person performing it.

This Essay argues that this fracture raises an underexamined civil rights problem. As artificial intelligence becomes a primary tool for legal preparation, access to protected legal reasoning increasingly depends on professional credential and platform tier rather than function. Because effective participation in legal processes—whether defending against criminal charges, seeking redress, or navigating administrative systems—requires the ability to prepare without compelled disclosure, the uneven protection of AI-assisted work risks entrenching structural inequality across the legal system. Individuals without access to counsel or enterprise-grade tools may be uniquely exposed, even as they rely most heavily on AI to exercise legal rights.

This Essay contends that this emerging regime implicates core concerns of civil rights law, including equal protection, procedural fairness, and meaningful access to legal institutions. It proposes a functional rule: work product protection should attach to the nature of legal preparation in anticipation of proceedings, not the professional identity of the preparer or the technological channel through which the work is produced.

I. Introduction: The Fracture Becomes Visible

On February 10, 2026, two federal courts decided the same question and reached opposite conclusions.[1] The split is not a procedural anomaly. It is the first visible surface of a structural

[1] *See* United States v. Heppner, No. 25-cr-00503-JSR, 2026 WL 436479 (S.D.N.Y. Feb. 17, 2026); Warner v. Gilbarco, Inc., No. 2:24-cv-12333, 2026 WL 373043 (E.D. Mich. Feb. 10, 2026). Judge Rakoff's bench

inequality forming inside AI-mediated legal systems. The question was whether materials a person created with the assistance of a publicly available AI platform—legal research, case analysis, strategic planning—are protected work product, or whether they are discoverable because the AI platform is not an attorney and its terms of service disclaim confidentiality. The answer, it turns out, depends on who the person is.

In *United States v. Heppner*, the U.S. District Court for the Southern District of New York held that a criminal defendant's AI-generated legal documents were not protected work product.[2] Bradley Heppner, facing securities fraud charges, had used a consumer AI platform to create thirty-one documents containing defense strategies, factual analyses, and legal arguments. He did so without his attorney's direction. The court, applying what it called the "core" of the work-product doctrine—"shelter[ing] the mental processes of the attorney, providing a privileged area within which he can analyze and prepare his client's case"—found that materials created by a non-lawyer without attorney involvement fell outside the doctrine's shelter.[3] The AI's terms of service, which permitted data collection and disclosure to third parties, reinforced the conclusion that Heppner had no reasonable expectation of confidentiality. The documents were deemed fully discoverable.

On the same day, in the Eastern District of Michigan, *Warner v. Gilbarco, Inc.* reached the opposite result.[4] Sohyon Warner, a pro se plaintiff in an employment discrimination case, had used consumer AI to prepare her case. When the defendants moved to compel production of all AI-related materials, the court denied the motion. Pro se litigants, it reasoned, occupy a unique dual role—"simultaneously the party and the advocate."[5] Their AI-assisted materials reflect their own mental impressions as the person litigating, and compelling production "would nullify work-product protection in nearly every modern drafting environment, a result no court has endorsed."[6] The AI was a tool, not a person; disclosure to a tool did not waive protection.

Within weeks, *Morgan v. V2X, Inc.* in the District of Colorado extended *Warner*'s reasoning, explicitly distinguishing *Heppner* on the ground that the criminal defendant there had acted independently of counsel while the pro se civil litigant was herself the advocate.[7] A doctrinal fracture has opened.

The emerging divide is not about whether AI is reliable, whether its outputs are accurate, or whether its use constitutes the unauthorized practice of law. It is about whether the

ruling was issued February 10, 2026; the written opinion followed February 17, 2026.

[2] Heppner, 2026 WL 436479.

[3] *Id.* at *3 (quoting United States v. Nobles, 422 U.S. 225, 238 (1975)).

[4] Warner, 2026 WL 373043.

[5] *Id.* at *3.

[6] *Id.* at *4.

[7] Morgan v. V2X, Inc., No. 25-cv-01991-SKC-MDB, 2026 WL 864223 (D. Colo. Mar. 30, 2026).

protection of legal preparation—the ability to research, analyze, and strategize without surrendering that work to the opposing side—attaches to the nature of the work or the status of the person performing it. This Essay enters an empty field. Existing scholarship treats AI as a civil rights threat in employment, housing, and criminal sentencing contexts, but has not yet examined how AI's role in legal preparation itself implicates equal protection, procedural fairness, and access to courts. This Essay argues that the fracture reveals a structural inequality in the distribution of protected legal thought, one that AI has made visible and that civil rights law is uniquely positioned to address. At its most extreme, this fracture leaves crime victims—who are not parties to the criminal proceeding, not represented by the prosecutor, and not entitled to appointed counsel—with no work-product framework for their preparation at all. The core thesis is that work-product protection is increasingly determined by professional credential and subscription tier rather than by function, and that this development conditions the exercise of fundamental rights on status and access.

The argument proceeds in six parts. Part II describes how AI has transformed access to legal capability without a corresponding expansion of protection. Part III examines the constitutional infrastructure—the technical, contractual, and institutional systems—that now mediates the protection of legal thought. Part IV anchors the fracture in civil rights law: equal protection, procedural fairness, and access to courts. Part V proposes a functional rule that would align protection with preparation rather than status. Part VI identifies materials—including a recently discovered conservation law—that make operational verification of protected preparation increasingly feasible. A brief conclusion returns to the civil rights stakes.

II. The Transformation of Legal Capability: Capability Without Protection

Work-product doctrine rests on a factual predicate that AI has destroyed. The doctrine, recognized in *Hickman v. Taylor* and codified in Federal Rule of Civil Procedure 26(b)(3), protects "documents and tangible things that are prepared in anticipation of litigation or for trial by or for another party or its representative."[8] The historical justification for tying protection to attorney status was straightforward: only lawyers could generate the kind of legal analysis the doctrine was designed to shelter. The ability to perform attorney-quality legal work tracked the professional credential. Extending protection to the attorney's agents—paralegals, investigators, consultants—was a practical extension, not a theoretical challenge, because those agents operated under the attorney's direction. The system was credential-gated because the capability was credential-gated.

[8] 329 U.S. 495 (1947); Fed. R. Civ. P. 26(b)(3)(A). The criminal equivalent in New York is CPL § 245.65.

AI has broken that link. A non-lawyer with access to a publicly available large language model can now research case law, analyze statutory elements, map evidence to legal standards, identify procedural requirements, evaluate defense strategies, and draft pleadings at a level of sophistication that would have required a trained attorney a decade ago.[9] This is not a speculative claim about future technology. It is the present reality for millions of individuals who navigate the legal system without counsel—pro se civil litigants, criminal defendants supplementing appointed counsel's work, administrative claimants, and crime victims preparing to engage with prosecutors.[10]

The collapse of the credential-to-capability link creates a civil rights gap. Effective participation in legal proceedings requires preparation, and preparation under threat of compelled disclosure is not meaningful preparation. If a person researching their legal rights knows that every query, every analysis, every drafted strategy can be demanded by the opposing side—while the opposing side's lawyer prepares in the shelter of work-product protection—the playing field is structurally tilted. The tilt does not arise from differences in the quality of the work. It arises from the asymmetry of protection.

The *Heppner/Warner* split makes this asymmetry concrete. The criminal defendant who used AI independently of counsel lost all protection. The pro se civil plaintiff who used the same kind of AI for the same kind of preparation retained protection. The decisive variable was not the nature of the work, the content of the prompts, or the quality of the output. It was the user's status relative to the attorney-client relationship.[11] For represented parties acting outside that relationship, AI-assisted work is exposed. For self-represented parties, AI-assisted work may be sheltered. For crime victims—who are not parties, not represented by the prosecutor, and not self-represented litigants—no framework attaches at all.[12]

The doctrinal distinction between these categories has a functional consequence: people who cannot afford counsel or who fall through the cracks of the existing protection categories are the people most likely to rely on AI for legal preparation, and their preparation is the least protected. AI has expanded access to legal capability without expanding access to protected legal preparation. The capability is distributed; the protection is not.

---

[9] *See generally* Danielle Keats Citron & Frank Pasquale, *The Scored Society: Due Process for Automated Predictions*, 89 WASH. L. REV. 1 (2014) (discussing the shift in analytical capability from human professionals to automated systems).

[10] The gap in legal representation is well documented. *See* Legal Servs. Corp., *The Justice Gap: The Unmet Civil Legal Needs of Low-Income Americans* (2022) (finding that 86% of civil legal problems reported by low-income Americans received inadequate or no legal help).

[11] *Compare* Heppner, 2026 WL 436479, at *4 (no counsel direction, no protection) *with* Warner, 2026 WL 373043, at *3 (self-represented party is own advocate, protection applies).

[12] Crime victims are not parties to the criminal proceeding; the prosecutor represents the state. *See* Linda R.S. v. Richard D., 410 U.S. 614, 619 (1973). Their statutory rights under 18 U.S.C. § 3771 do not include the right to appointed counsel or to work-product protection for their own preparation materials.

III. Constitutional Infrastructure and the Privatization of Protected Thought

The fracture is compounded by the technological and contractual architecture through which AI-assisted legal preparation is now delivered. The infrastructure matters because it is the mechanism through which protection is distributed.

A. The Wrapper Ecosystem

The leading AI tools marketed specifically to the legal profession are built on the same foundation models available to the public. Harvey, an AI platform used by major law firms, integrates Anthropic's Claude alongside OpenAI's GPT models.[13] CoCounsel, Thomson Reuters's AI assistant, is built on OpenAI's GPT-4 and holds a contract with the federal judiciary.[14] Lexis+ AI and Protégé draw on multiple foundation models. These enterprise platforms are sold through subscription tiers that range from tens of thousands to hundreds of thousands of dollars annually. They are marketed with the promise that materials created within their environments are protected as attorney work product.[15]

Consumer versions of the same underlying models—Claude, ChatGPT, Gemini—are available for free or at modest monthly cost. They produce outputs of materially comparable analytical quality. But the legal profession's discourse, reinforced by *Heppner*'s reasoning, treats materials created on consumer platforms by non-lawyers as presumptively unprotected.[16]

Publicly available information suggests that the distinction between enterprise and consumer tools lies primarily in contractual terms, data handling policies, and integration layers, not in demonstrable differences in the underlying analytical models.[17] Enterprise agreements typi-

[13] *See* Proskauer Rose LLP, *SDNY Addresses Privilege and Work Product Implications of Using Unsecured Public AI Tools* (Feb. 18, 2026), https://www.proskauer.com/alert/sdny-addresses-privilege-and-work-product-implications-of-using-unsecured-public-ai-tools [https://perma.cc/4JWU-BPWQ] (confirming Heppner's use of Anthropic's Claude); Harvey, *Opus 4.6, Now Live in Harvey* (2026), https://www.harvey.ai/blog/opus-4-6-now-live-in-harvey [https://perma.cc/7JGW-EWA5] (announcing integration of Anthropic's Claude Opus into the Harvey platform); *see also* Anthropic, *Customer Story: Harvey*, https://claude.com/customers/harvey [https://perma.cc/E2R7-98JK] (describing Harvey's use of Claude).

[14] *See* Press Release, Thomson Reuters, *Westlaw Precision with CoCounsel, Practical Law, and CoCounsel to be provided to US Federal Courts as the essential information provider for the Federal Judiciary* (Apr. 9, 2025), https://www.thomsonreuters.com/en/press-releases/2025/april/westlaw-precision-with-cocounsel-practical-law-and-cocounsel-to-be-provided-to-us-federal-courts-as-the-essential-information-provider-for-the-federal-judiciary [https://perma.cc/Q5EU-KE4L] (announcing multi-year contract with the Administrative Office of the U.S. Courts).

[15] *See, e.g.*, LexisNexis, *Lexis+ with Protégé* (formerly Lexis+ AI), https://www.lexisnexis.com/en-us/products/lexis-plus-ai.page [https://perma.cc/J9K6-SWYZ] (describing enterprise-grade security and confidentiality for legal work).

[16] *See* Heppner, 2026 WL 436479, at *4 (citing the consumer AI platform's terms of service as a basis for finding no reasonable expectation of confidentiality).

[17] *See generally* Legaltech Hub, *Generative AI Toolkit*, https://www.legaltechnologyhub.com/topics/process-

cally prohibit the model provider from training on customer data, include SOC 2 compliance certifications, and add confidentiality provisions. They do not, however, publicly establish that the model's internal computation is segregated between enterprise and consumer data paths. Publicly available materials do not establish that enterprise and consumer processing paths are architecturally segregated. The separation is addressed through data governance promises, not architectural isolation.[18]

As a result, the distinction between protected and unprotected AI-assisted legal preparation appears, in practice and as emerging cases suggest, to track access tier and user status rather than demonstrable differences in underlying analytical capability. This is not a technology distinction dressed in legal clothing. It is a credentialing product that signals: *this is attorney work, accessed through attorney channels, therefore the output is protected.* The same model, accessed through the consumer platform by a non-attorney, produces output that is not protected—not because the work is different, but because the access tier identifies the user as a non-attorney.

B. Power Redistribution

This architecture directly answers a question the Stanford Law Review collection poses: how do technological systems redistribute power between private actors and the state?[19] The answer, in this context, is that they redistribute who gets protected legal thought.

Enterprise legal AI platforms have taken publicly available model capability and wrapped it in a contractual and credentialing layer that determines whether the constitutional protection of work product attaches. Courts have not consciously decided to create a two-tier system. They have, however, applied a doctrine designed for the pre-AI era in a manner that effectively delegates the protection of legal reasoning to private subscription architecture. The result is a system in which protected legal preparation flows through channel partners, not through constitutional entitlement.

When courts, privilege doctrine, and procedural rules ratify an arrangement in which protected legal reasoning is functionally available only through licensed intermediaries and their enterprise tools, the inequality is state-backed. Private contracts are the mechanism of the inequality, not a defense to it. The delegation of constitutional protection into private channel

improvement/generative-ai-toolkit/ [https://perma.cc/98ER-ZTT3]; *see also* UC Davis Mabie Law Library, *Generative AI Tools and Resources for Law Students*, https://libguides.law.ucdavis.edu/genaiforlawstudents/tools [https://perma.cc/GRE9-VP5E] (citing Legaltech Hub's March 2026 mapping of 1,014 generative AI product placements in legal tech).

[18] Publicly available terms of service, privacy policies, and industry descriptions reviewed as of April 2026 do not identify architectural segregation commitments between enterprise and consumer processing paths.

[19] *See* Stanford Law Review Online, Call for Submissions: *Technology, Artificial Intelligence, and the Future of Civil Rights* (2026), https://www.stanfordlawreview.org/online/ [https://perma.cc/AYZ4-YXSJ].

architecture raises questions that extend beyond privilege doctrine into the domain of civil rights.

IV. Civil Rights at the Fault Line: Equality, Fairness, and Access

The AI work-product fracture implicates three foundational civil rights principles: equal protection, procedural fairness, and access to courts.

A. Equal Protection and Status-Based Distinction

Two persons preparing for litigation—one a lawyer through Harvey, the other a self-represented litigant through consumer AI—may produce materially comparable legal analysis. The lawyer's work is protected. The self-represented litigant's work, under *Heppner*'s logic, may not be. The distinction does not rest on the function the work serves or the care with which it was prepared. It rests on the professional status of the person who prepared it and the platform tier through which it was produced.

This produces a status-based distinction in the protection of legal preparation that implicates equal protection principles in ways courts are only beginning to confront. *Morgan* itself recognized that conditioning AI work-product protection on attorney involvement "would further disadvantage unrepresented litigants"—the precise distributive consequence that civil rights scrutiny exists to address.[20] The historical justification for credential-based asymmetry—that lawyers had unique analytical capacity—collapses when the same AI model generates materially comparable analysis for both. When function detaches from credential, a doctrine that ties protection to credential tracks status, not function. That classification, when it burdens the population least able to afford the alternative, is one that civil rights law has reason to scrutinize.

Work-product protection is not merely a procedural convenience. It is a precondition to the exercise of legal rights. A person who cannot research, analyze, and strategize in confidence cannot meaningfully prepare a defense, prosecute a claim, or navigate an administrative proceeding. In this sense, work-product protection stands in the same structural relationship to adversarial legal participation that access to transcripts stands to the right of appeal in *Griffin v. Illinois*, that access to counsel stands to the right of defense in *Gideon v. Wainwright*, and that the right of self-representation stands to the right to present one's own case in *Faretta v. California*. Each of those doctrines recognizes that the right itself is hollow

[20] *Cf.* Griffin v. Illinois, 351 U.S. 12, 19 (1956) ("There can be no equal justice where the kind of trial a man gets depends on the amount of money he has."); *see also* M.L.B. v. S.L.J., 519 U.S. 102 (1996) (extending Griffin's principle to civil proceedings involving fundamental interests); Douglas v. California, 372 U.S. 353 (1963) (right to counsel on first appeal for indigent defendants).

without the conditions that make its exercise possible. A doctrine that conditions the ability to prepare—to think legally in private—is functionally part of the right itself, not an adjunct to it. Griffin's principle—that "there can be no equal justice where the kind of trial a man gets depends on the amount of money he has"—applies directly when what determines that access is subscription tier rather than financial means per se. The inequality is the same. Only the mechanism has changed.[21]

B. Procedural Fairness and Technological Due Process

The collection's framing of procedural fairness is directly implicated. Preparation exposure is procedural disadvantage. When one party prepares under shelter—knowing their research, analysis, and strategy cannot be compelled—and another party prepares under threat of discovery for performing the identical function, the proceeding is not procedurally fair. The disadvantage is not in the quality of the analysis; it is in the safety of the process.[22]

This is a form of the technological due process concern that scholars have identified in other AI-mediated contexts: systems are producing procedural outcomes that doctrine has not yet recognized as fairness problems.[23] The AI work-product fracture fits that pattern. The infrastructure through which legal preparation is conducted is distributing protection unevenly, and courts have not yet confronted the procedural unfairness that results.

C. Access to Courts and the Right to Petition

Preparation is a prerequisite to participation. The right to petition the government for redress of grievances,[24] the right of self-representation recognized in *Faretta v. California*,[25] and the due process right to a fair proceeding[26] all require the ability to prepare. If the protection regime penalizes use of the only affordable preparation tool, access to courts is effectively conditioned on the ability to afford enterprise-tier legal infrastructure or to retain counsel.

This burden falls most heavily on populations that are already structurally underserved: indigent criminal defendants using AI before counsel is appointed or to supplement appointed counsel's work;[27] pro se civil litigants in housing, employment, and benefits cases; administra-

[21] *See* Gideon v. Wainwright, 372 U.S. 335 (1963); Faretta v. California, 422 U.S. 806 (1975); Griffin, 351 U.S. at 19. *See also* Bounds v. Smith, 430 U.S. 817 (1977) (prisoners have a fundamental right of access to the courts); Lewis v. Casey, 518 U.S. 343 (1996) (elaborating on the access-to-courts right and its limitations).

[22] *See* Mathews v. Eldridge, 424 U.S. 319, 335 (1976) (procedural due process requires a balancing of the private interest, the risk of erroneous deprivation, and the government's interest).

[23] *See* Danielle Keats Citron, *Technological Due Process*, 85 WASH. U. L. REV. 1249 (2008); Cary Coglianese & David Lehr, *Regulating by Robot: Administrative Decision Making in the Machine-Learning Era*, 105 GEO. L.J. 1147 (2017).

[24] U.S. CONST. amend. I.

[25] 422 U.S. 806 (1975).

[26] U.S. CONST. amend. XIV, § 1.

[27] *See* Clio, *Does Using AI Waive Attorney-Client Privilege? Breaking Down the United States v. Heppner*

tive claimants navigating immigration, Social Security disability, and licensing proceedings; and crime victims who are not parties to the criminal case, not represented by the prosecutor, and not entitled to appointed counsel.[28] For each, AI is the primary tool for legal preparation. For each, the protection of that preparation is either absent or ambiguous under current doctrine.

The fracture extends into every domain where individuals must prepare to exercise legal rights without the shelter of counsel. Civil rights law has addressed analogous structural asymmetries—discriminatory hiring tools, automated benefits denials, proxy discrimination in lending—by identifying the ways technically neutral rules produce systematically unequal outcomes.[29] The AI work-product fracture is the same species of problem.

V. A Rule for the Age of AI: Functional Protection

The solution is not to abandon work-product protection or to extend it indiscriminately. It is to align the protection with the function it serves: sheltering the preparation that enables meaningful participation in legal proceedings.

The proposed rule is functional. Work-product protection should attach to the nature of legal preparation in anticipation of litigation or legal proceedings, not to the professional identity of the preparer or the technological channel through which the work is produced. AI-assisted material retains protection when prepared in anticipation of proceedings, unless the proponent demonstrates actual adversarially relevant exposure—beyond ordinary computational processing.[30] This approach builds on decisions that predate the current AI era, including *Shih v. Petal Card, Inc.*, which held that work-product protection extends to litigation preparation by a non-attorney corporate representative acting in anticipation of suit.[31]

The doctrine that work-product protection incentivizes the attorney-client relationship serves

Case, CLIO BLOG (2026), https://www.clio.com/blog/ai-attorney-client-privilege-heppner/ [https://perma.cc/GE4K-94G4] (analyzing privilege risks where defendants use consumer AI without counsel direction).

[28] *See supra* note 12.

[29] *See generally* Solon Barocas & Andrew D. Selbst, *Big Data's Disparate Impact*, 104 CALIF. L. REV. 671 (2016) (arguing that facially neutral data processing can produce discriminatory outcomes); Anya E.R. Prince & Daniel Schwarcz, *Proxy Discrimination in the Age of Artificial Intelligence and Big Data*, 105 IOWA L. REV. 1257 (2020).

[30] Courts have begun drawing distinctions between technical processing and legally meaningful disclosure in other digital contexts. *Cf.* Carpenter v. United States, 138 S. Ct. 2206, 2219 (2018) (quoting Riley v. California, 134 S. Ct. 2473, 2488 (2014)) (holding that diminished privacy interests in information shared with a third party do not mean "the Fourth Amendment falls out of the picture entirely").

[31] 565 F. Supp. 3d 557, 574 (S.D.N.Y. 2021) (extending work-product protection to litigation preparation by a non-attorney corporate representative acting in anticipation of suit).

important values—competence, ethics regulation, professional accountability. However, that incentive argument has no force in the circumstances this rule addresses: where counsel is unavailable or unaffordable, and the alternative is no preparation at all. The functional rule protects preparation in the gap, not in place of counsel.

A deeper objection challenges the rule on doctrinal foundations: AI output, the argument runs, lacks the human mental impressions that *Hickman* and Rule 26(b)(3) exist to protect.[32] The objection is misplaced. The functional rule does not protect AI output as algorithmic product; it protects the human preparation activity—the prompts, the directional choices, the integration of output into a litigation strategy—that embeds the preparer's mental impressions as surely as a paralegal's research memorandum does. The copyright analogy under *Thaler v. Perlmutter* is distinguishable on this same ground: copyright requires a human author whose creative expression resides in the work, while work product requires a human preparer whose strategic preparation is reflected in the materials. The AI is the substrate; the legal actor remains the human.

This rule requires courts to draw a distinction that *Heppner* collapsed. The mere fact that an AI system processes a prompt does not necessarily constitute disclosure to an adversary or third party in the traditional sense, particularly where no evidence of retention or adversarial access is shown.[33] The legally relevant events are onward sharing, human review by platform staff, government subpoena of stored data, or actual adversarial access—not the computational inference itself. Treating the potential for future training or data retention as an immediate waiver of protection is equivalent to saying that an email provider's ability to scan messages for spam waives attorney-client privilege over the scanned communications. The doctrine does not require that result, and courts should not adopt it.

The functional rule is consistent with the text of Federal Rule of Civil Procedure 26(b)(3)(A), which protects materials prepared "by or for another party or its representative" without credential-limiting language.[34] *Warner* and *Morgan* already take this path. Extending it to all persons preparing in anticipation of litigation—regardless of whether they are represented, self-represented, or preparing to participate as a crime victim in a criminal proceeding—is a doctrinal development, not a statutory amendment.

Three paths are available. Judicial reinterpretation can adopt the functional rule as *Warner*

[32] *See* Hickman v. Taylor, 329 U.S. 495, 510–11 (1947); United States v. Nobles, 422 U.S. 225, 238 (1975); *see also* Roberts v. Americable Int'l, Inc., 883 F. Supp. 499 (E.D. Cal. 1995) (denying work-product protection where materials revealed only the conversants' thought processes, not the attorney's); Thaler v. Perlmutter, 130 F.4th 1039 (D.C. Cir. 2025) (holding AI cannot be a copyright author because authorship requires a human creator).

[33] *See* Warner, 2026 WL 373043, at *4 ("Generative AI programs are tools, not persons.").

[34] Fed. R. Civ. P. 26(b)(3)(A). The rule's text covers "a party or its representative," which on its face does not limit the representative to an attorney. *See also Shih*, 565 F. Supp. 3d at 574 (extending protection to a non-attorney representative).

and *Morgan* have begun to do. Rulemaking can amend the Federal Rules of Civil Procedure and parallel state rules to explicitly extend protection to AI-assisted preparation by self-represented parties and by parties preparing for any legal proceeding, regardless of platform tier. And legislation can address the specific gaps—for example, targeted amendments to the Crime Victims' Rights Act[35] or model provisions for state victims' rights campaigns—to ensure that crime victims' AI-assisted preparation materials are not discoverable by the very defendants whose conduct they are documenting.[36]

Each path is a civil rights intervention. Each recognizes that the protection of legal preparation is not a professional courtesy but a prerequisite to the equal exercise of legal rights.

VI. Materials for Building: The Infrastructure Beneath the Rule

A doctrinal rule is necessary but not itself sufficient. To be operational, a rule that protects AI-assisted legal preparation requires the capacity to verify that the protection has been preserved—that the preparation was conducted in confidence and that its protected character survived the technological pipeline. The infrastructure for that verification is beginning to emerge.

One material is a recently discovered conservation law describing how meaning behaves under transformation.[37] The Conservation Law of Commitment establishes that under governed conditions, the commitment kernel of a signal—its obligations, prohibitions, permissions, and constraints—is conserved; under ungoverned conditions, it decays in predictable, measurable ways. The law is substrate-agnostic. It applies to human transformation, AI transformation, and any hybrid.

Applying that framework to the work-product context—defining the kernel as "sheltered preparation in anticipation of proceedings" and implementing governed transformation protocols—would make the functional rule's protection demonstrable, not merely doctrinal. A person preparing for litigation could demonstrate not only that they intended their work to be confidential but that the confidentiality was structurally preserved—through cryptographic lineage, fidelity verification gates, and recursion testing.[38]

[35] 18 U.S.C. § 3771. The Act currently provides victims with the right to be heard, the right to notice, and the right to restitution, but does not address the protection of their litigation preparation materials.

[36] *See* Marsy's Law for All, *Model Amendment Language*, https://www.marsyslaw.us [https://perma.cc/4G8H-2WWF]. The model amendment does not include provisions for AI-assisted preparation or work-product protection.

[37] [Author], *A Conservation Law for Commitment in Language Under Transformative Compression and Recursive Application* (Mar. 19, 2026) (unpublished preprint), https://doi.org/10.5281/zenodo.18792459.

[38] *See* [Author], *supra* note 36, at §§ 4–5 (describing the six-gate verification protocol and recursion testing methodology).

These are not speculative technologies. They are operational components of a governance architecture that has been publicly described and deposited.[39] They do not require opening proprietary models or revealing trade secrets. They certify the result, not the internals. They allow courts, agencies, and parties to ask the question that *Heppner* and *Warner* each approached obliquely: did the protection survive the pipeline?

This Essay does not ask courts to adopt any particular technical architecture. It argues only that verifiable preservation of confidentiality and legal meaning is becoming possible—and that the functional rule proposed here is designed to take advantage of that possibility as it matures.

VII. Conclusion: Whose Legal Thought Stays Protected?

The civil rights project has always been about identifying the ways law formally grants equality while practically permitting inequality to persist through structural arrangements. The AI work-product fracture is one of those arrangements. Courts have not consciously decided to create a two-tier system of protected legal thought. They have done so inadvertently, by applying a doctrine designed for the pre-AI world to a world in which the factual predicate of that doctrine—that legal analytical capability tracks professional credential—has collapsed.

The result is a system in which lawyers prepare under shelter, enterprise subscribers prepare under shelter, and everyone else prepares under threat. The threat falls most heavily on the populations that civil rights law exists to protect.

Artificial intelligence is reshaping how individuals prepare to assert and defend their rights. Civil rights law is at an inflection point—courts are narrowing antidiscrimination doctrine, rethinking protected classes, and grappling with technologically mediated harm in every domain. The work-product fracture belongs in that conversation. As AI becomes the primary tool through which unrepresented individuals access the legal system, civil rights law must confront a new question: not only who is protected, but whose legal thought is allowed to remain protected at all.

Footnotes

---

[39] [Author], *Experimental Record for A Conservation Law for Commitment in Language Under Transformative Compression and Recursive Application* (EXP-001 to EXP-007), https://doi.org/10.5281/zenodo.19105225.

Note: Footnotes 36–38 contain self-citations anonymized for blind review. Author identifying information provided separately in submission form fields.